

## 10.1 卷积神经网络

### layers.Conv2D

```
In [1]: import tensorflow as tf
In [2]: from tensorflow.keras import layers

In [6]: layer=layers.Conv2D(4, kernel_size=5,strides=1,padding='valid')
In [8]: out=layer(x)
Out[9]: TensorShape([1, 28, 28, 4])

In [10]: layer=layers.Conv2D(4, kernel_size=5,strides=1,padding='same')
In [11]: out=layer(x)
Out[12]: TensorShape([1, 32, 32, 4])

In [13]: layer=layers.Conv2D(4, kernel_size=5,strides=2,padding='same')
In [14]: out=layer(x)
Out[15]: TensorShape([1, 16, 16, 4])

In [16]: layer.call(x).shape
Out[16]: TensorShape([1, 16, 16, 4])
```

### weight & bias

```
In [13]: layer=layers.Conv2D(4, kernel_size=5,strides=2,padding='same')
In [14]: out=layer(x)
Out[15]: TensorShape([1, 16, 16, 4])

In [17]: layer.kernel
<tf.Variable 'conv2d_3/kernel:0' shape=(5, 5, 3, 4) dtype=float32, numpy=
array([[[[-0.16160963,  0.04107726, -0.09828208, -0.00601757],
          [-0.02003701,  0.01415607, -0.07604317, -0.12557343],
          [-0.11157566,  0.1328298 ,  0.14624669, -0.04775226]], ...

In [18]: layer.bias
Out[18]: <tf.Variable 'conv2d_3/bias:0' shape=(4,) dtype=float32, numpy=array([0.,
0., 0., 0.], dtype=float32)>
```

### nn.conv2d

```
In [21]: w=tf.random.normal([5,5,3,4])
In [22]: b=tf.zeros([4])
In [23]: x.shape
Out[23]: TensorShape([1, 32, 32, 3])

In [29]: out=tf.nn.conv2d(x, w, strides=1, padding='VALID')
Out[30]: TensorShape([1, 28, 28, 4])

In [31]: out = out + b
Out[32]: TensorShape([1, 28, 28, 4])

In [33]: out=tf.nn.conv2d(x,w,strides=2,padding='VALID')
Out[34]: TensorShape([1, 14, 14, 4])
```

## 卷积梯度的求导

$x_{00}$	$x_{01}$	$x_{02}$
$x_{10}$	$x_{11}$	$x_{12}$
$x_{20}$	$x_{21}$	$x_{22}$

$w_{00}$	$w_{01}$
$w_{10}$	$w_{11}$

$o_{00}$	$o_{01}$
$o_{10}$	$o_{11}$

0
1
0
0

$$o_{00} = x_{00} * w_{00} + x_{01} * w_{01} + x_{10} * w_{10} + x_{11} * w_{11} + b$$

$$o_{01} = x_{01} * w_{00} + x_{02} * w_{01} + x_{11} * w_{10} + x_{12} * w_{11} + b$$

$$o_{10} = x_{10} * w_{00} + x_{11} * w_{01} + x_{20} * w_{10} + x_{21} * w_{11} + b$$

$$o_{11} = x_{11} * w_{00} + x_{12} * w_{01} + x_{21} * w_{10} + x_{22} * w_{11} + b$$

$$\frac{\partial Loss}{\partial w_{00}} = \sum_{i \in \{00, 01, 10, 11\}} \frac{\partial Loss}{\partial o_i} \frac{\partial o_i}{\partial w_{00}}$$


$$\frac{\partial o_{00}}{\partial w_{00}} = \frac{\partial (x_{00} * w_{00} + x_{01} * w_{01} + x_{10} * w_{10} + x_{11} * w_{11} + b)}{\partial w_{00}} = x_{00}$$

## 10.2 池化与采样

- ✓ Pooling       $\leftrightarrow$  下采样
- ✓ upsample     $\leftrightarrow$  上采样
- ✓ ReLU

### 下采样

- stride=2



```

In [36]: x # TensorShape([1, 14, 14, 4])

In [37]: pool=layers.MaxPool2D(2,strides=2)
In [38]: out=pool(x)
Out[39]: TensorShape([1, 7, 7, 4])

In [40]: pool=layers.MaxPool2D(3,strides=2)
In [41]: out=pool(x)
Out[42]: TensorShape([1, 6, 6, 4])

In [44]: out=tf.nn.max_pool2d(x, 2,strides=2,padding='VALID')
Out[45]: TensorShape([1, 7, 7, 4])
    
```

## 补充，关于卷积核后的图像尺寸计算，

### 1、卷积后尺寸的计算

$out\_height = (in\_height + 2 * padding - filter\_height) / strides[1] + 1$

$out\_width = (in\_width + 2 * padding - filter\_width) / strides[2] + 1$

### 2、卷积参数 *same* 和 *valid* 运算之后的维度计算

#### (1) same

$out\_height = \lceil \text{float}(in\_height) \rceil / \text{float}(strides[1])$

$out\_width = \lceil \text{float}(in\_width) \rceil / \text{float}(strides[2])$

#### (2) valid

$out\_height = \lceil \text{float}(in\_height - filter\_height + 1) \rceil / \text{float}(strides[1])$

$out\_width = \lceil \text{float}(in\_width - filter\_width + 1) \rceil / \text{float}(strides[2])$

关于参数：

padding: SAME 和 VALID 两种形式

filter: [5,5,1,32] 表示 5\*5 的卷积核，1 个 channel，32 个卷积核。

strides: [1,4,4,1] 表示横向和竖向的步长都为 4

## 上采样

```
In [47]: x=tf.random.normal([1,7,7,4])
In [48]: layer=layers.UpSampling2D(size=3)
In [49]: out=layer(x)
Out[50]: TensorShape([1, 21, 21, 4])

In [51]: layer=layers.UpSampling2D(size=2)
In [52]: out=layer(x)
Out[53]: TensorShape([1, 14, 14, 4])
```

## ReLU

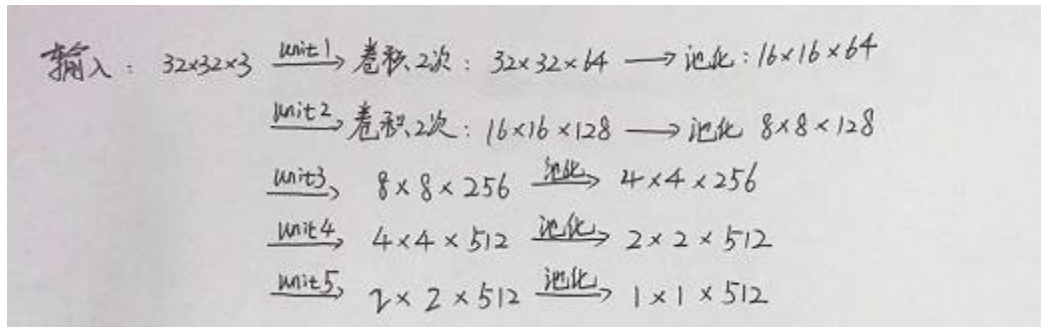

```
In [55]: x=tf.random.normal([2,3])
<tf.Tensor: id=154, shape=(2, 3), dtype=float32, numpy=
array([[ -1.533682,  -2.705335,  0.36354962],
       [ 0.00713745,  0.69756126,  0.8053344 ]], dtype=float32)>

In [57]: tf.nn.relu(x)
<tf.Tensor: id=156, shape=(2, 3), dtype=float32, numpy=
array([[0., 0., 0.36354962],
       [0.00713745, 0.69756126, 0.8053344 ]], dtype=float32)>

In [59]: layers.ReLU()(x)
<tf.Tensor: id=158, shape=(2, 3), dtype=float32, numpy=
array([[0., 0., 0.36354962],
       [0.00713745, 0.69756126, 0.8053344 ]], dtype=float32)>
```

# 实战

## 网络结构



代码:

```
import tensorflow as tf
from tensorflow.keras import layers, optimizers, datasets, Sequential
tf.random.set_seed(2345)

# 5 unites of conv + maxpooling
conv_layers = [
    # unit 1
    layers.Conv2D(64, kernel_size=[3,3], padding='same', activation=tf.nn.relu),
    layers.Conv2D(64, kernel_size=[3,3], padding='same', activation=tf.nn.relu),
    layers.MaxPool2D(pool_size=[2,2], strides=2, padding='same'),

    # unit 2
    layers.Conv2D(128, kernel_size=[3,3], padding='same', activation=tf.nn.relu),
    layers.Conv2D(128, kernel_size=[3,3], padding='same', activation=tf.nn.relu),
    layers.MaxPool2D(pool_size=[2,2], strides=2, padding='same'),

    # unit 3
    layers.Conv2D(256, kernel_size=[3,3], padding='same', activation=tf.nn.relu),
    layers.Conv2D(256, kernel_size=[3,3], padding='same', activation=tf.nn.relu),
    layers.MaxPool2D(pool_size=[2,2], strides=2, padding='same'),

    # unit 4
    layers.Conv2D(512, kernel_size=[3,3], padding='same', activation=tf.nn.relu),
    layers.Conv2D(512, kernel_size=[3,3], padding='same', activation=tf.nn.relu),
    layers.MaxPool2D(pool_size=[2,2], strides=2, padding='same'),

    # unit 5
    layers.Conv2D(512, kernel_size=[3,3], padding='same', activation=tf.nn.relu),
    layers.Conv2D(512, kernel_size=[3,3], padding='same', activation=tf.nn.relu),
    layers.MaxPool2D(pool_size=[2,2], strides=2, padding='same'),

    # final layer
    layers.Dense(1000, activation='softmax')
]
```

```

layers.MaxPool2D(pool_size=[2,2], strides=2, padding='same'),

# unit 5
layers.Conv2D(512, kernel_size=[3,3], padding='same', activation=tf.nn.relu),
layers.Conv2D(512, kernel_size=[3,3], padding='same', activation=tf.nn.relu),
layers.MaxPool2D(pool_size=[2,2], strides=2, padding='same')
]

# 数据预处理与加载数据集
def preprocess(x, y):
    x = tf.cast(x, dtype=tf.float32) / 255.
    y = tf.cast(y, dtype=tf.int32)
    return x,y
(x,y),(x_test,y_test) = datasets.cifar100.load_data()
y = tf.squeeze(y, axis=1)
y_test = tf.squeeze(y_test, axis=1)

train_db = tf.data.Dataset.from_tensor_slices((x,y))
train_db = train_db.shuffle(1000).map(preprocess).batch(128)
test_db = tf.data.Dataset.from_tensor_slices((x,y))
test_db = test_db.map(preprocess).batch(64)

def main():
    # [b, 32, 32, 3] ==> [b, 1, 1, 512]
    conv_net = Sequential(conv_layers)
    # conv_net.build(input_shape=[None, 32, 32,3])

    # 全连接层
    fc_net = Sequential([
        layers.Dense(256, activation=tf.nn.relu),
        layers.Dense(128, activation=tf.nn.relu),
        layers.Dense(100, activation=None)
    ])
    # 将上面两个网络给 build 一下
    conv_net.build(input_shape=[None, 32, 32, 3])
    fc_net.build(input_shape=[None, 512])

    # 设置优化器
    optimizer = optimizers.Adam(lr=1e-4)
    # 设置梯度参数
    variables = conv_net.trainable_variables + fc_net.trainable_variables

```

```

for epoch in range(50):
    for step, (x,y) in enumerate(train_db):
        with tf.GradientTape() as tape:
            # [b, 32, 32, 3] => [b, 1, 1, 512]
            out = conv_net(x)
            # flatten to [b, 512]
            out = tf.reshape(out, [-1, 512])
            # [b, 512] ==> [b, 100]
            logits = fc_net(out)
            # 对 y 进行 one_hot 编码
            y_onehot = tf.one_hot(y, depth=100)
            # 计算损失
            loss = tf.losses.categorical_crossentropy(y_onehot, l
ogits, from_logits=True)
            loss = tf.reduce_mean(loss)

            grads = tape.gradient(loss, variables)
            optimizer.apply_gradients(zip(grads, variables))

        if step % 100 == 0:
            print(epoch, step, 'loss', float(loss))

    total_num = 0
    total_correct = 0
    for x, y in test_db:
        out = conv_net(x)
        out = tf.reshape(out, [-1, 512])
        logits = fc_net(out)
        prob = tf.nn.softmax(logits, axis=1)
        pred = tf.argmax(prob, axis=1)
        pred = tf.cast(pred, dtype=tf.int32)

        correct = tf.cast(tf.equal(pred, y), dtype=tf.int32)
        correct = tf.reduce_sum(correct)

        total_num += x.shape[0]
        total_correct += int(correct)

    acc = total_correct / total_num
    print(epoch, 'acc:', acc)

main()

```

## LeNet-5



卷积核为11x11

- [ImageNet Classification with Deep Convolutional Neural Networks](#), NIPS 2012

## VGGNet: ILSVRC 2014 2<sup>nd</sup> place 11层到19层网络

Network	A, A-LRN	B	C	D	E
Number of parameters	133	133	134	138	144

- 

[Very Deep Convolutional Networks for Large-Scale Image Recognition](#), ICLR 2015

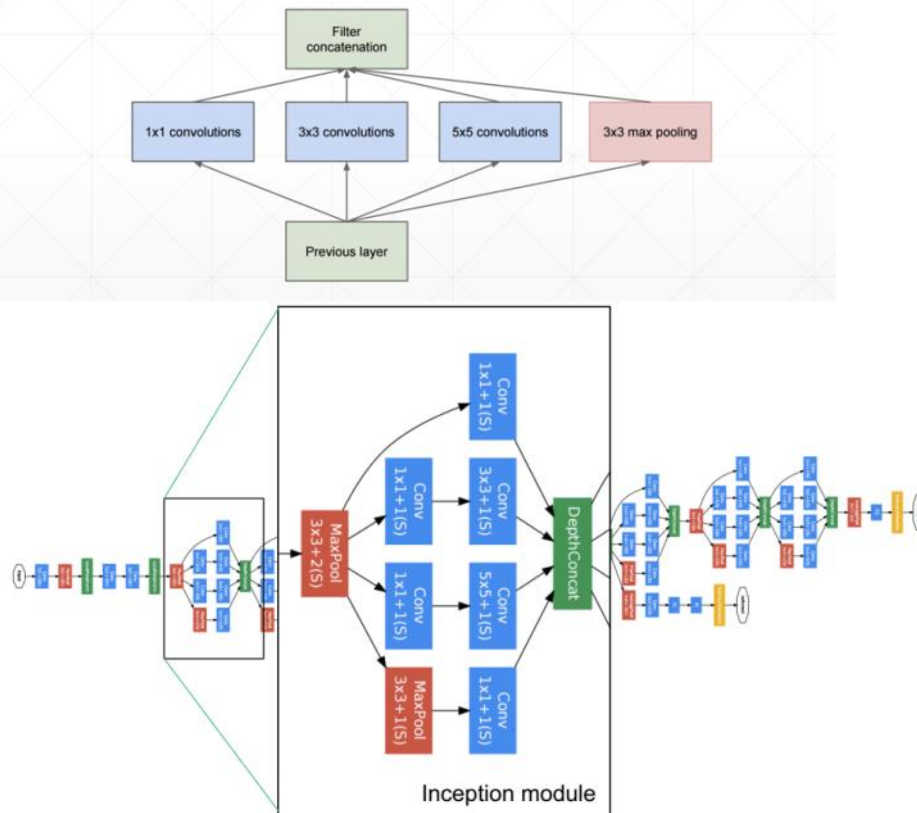


那么  $1 \times 1$  的卷积核有什么作用呢？

- less computation
- 可以改变图片的通道数，从而改变维度， $c_{in} \geq c_{out}$

## GooLeNet


- 1<sup>st</sup> in 2014 ILSVRC
- 22 layers



C. Szegedy et al., [Going deeper with convolutions](#), CVPR 2015

## ResNet 与 DenseNet


- The residual module
  - Introduce *skip* or *shortcut* connections (existing before in various forms in literature)
  - Make it easy for network layers to represent the identity mapping
  - For some reason, need to skip at least two layers



Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun,  
[Deep Residual Learning for Image Recognition](#), CVPR 2016 (Best Paper)



## Deeper residual module (bottleneck)




- Directly performing 3x3 convolutions with 256 feature maps at input and output:  
 $256 \times 256 \times 3 \times 3 \sim 600K$  operations
- Using 1x1 convolutions to reduce 256 to 64 feature maps, followed by 3x3 convolutions, followed by 1x1 convolutions to expand back to 256 maps:  
 $256 \times 64 \times 1 \times 1 \sim 16K$   
 $64 \times 64 \times 3 \times 3 \sim 36K$   
 $64 \times 256 \times 1 \times 1 \sim 16K$   
 Total:  $\sim 70K$

Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun,  
[Deep Residual Learning for Image Recognition](#), CVPR 2016 (Best Paper)

## 实现

### Basic Block



```
class BasicBlock(layers.Layer):
    def __init__(self, filter_num, stride=1):
        super(BasicBlock, self).__init__()

        self.conv1 = layers.Conv2D(filter_num, (3, 3), strides=stride, padding='same')
        self.bn1 = layers.BatchNorm2D(filter_num)
        self.relu = layers.Activation('relu')
        self.conv2 = layers.Conv2D(filter_num, (3, 3), strides=1, padding='same')
        self.bn2 = layers.BatchNorm2D(filter_num)

        if stride != 1:
            self.downsample = Sequential()
            self.downsample.add(layers.Conv2D(1, (1, 1), strides=stride))
            self.downsample.add(layers.BatchNorm2D(filter_num))
        else:
            self.downsample = lambda x: x

        self.stride = stride

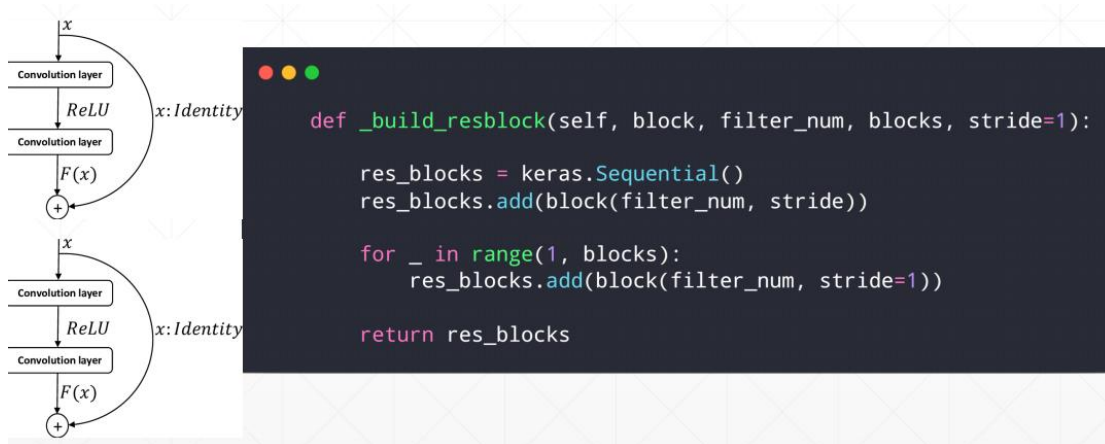
    def call(self, inputs, training=None):
        residual = self.downsample(inputs)

        conv1 = self.conv1(inputs)
        bn1 = self.bn1(conv1)
        relu1 = self.relu(bn1)
        conv2 = self.conv2(relu1)
        bn2 = self.bn2(conv2)

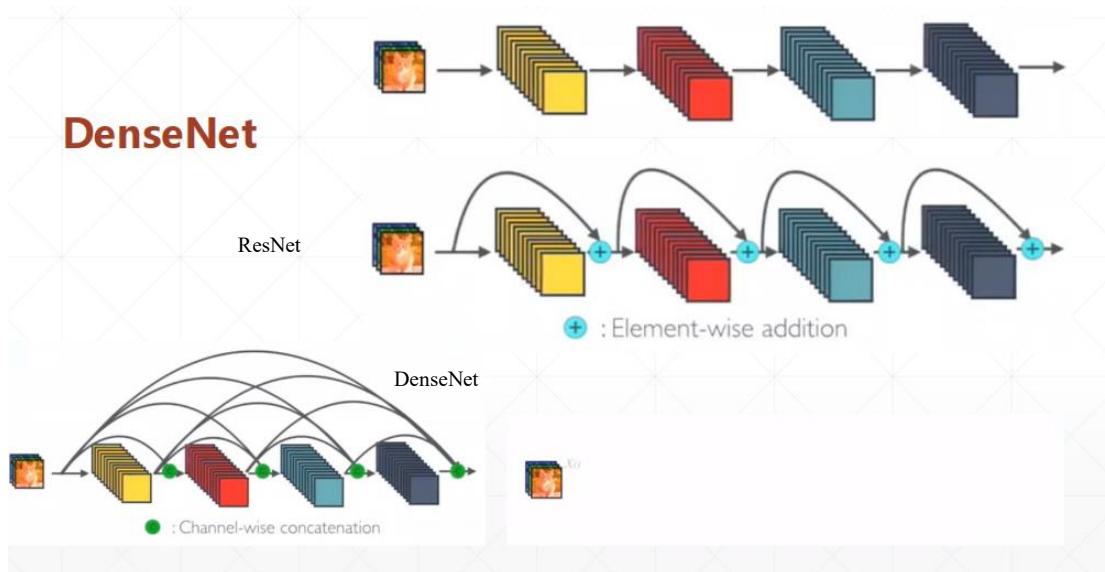
        add = layers.add([bn2, residual])
        out = self.relu(add)
        return out
```

### ResBlock

将 BasicBlock 堆叠起来，形成 ResBlock。

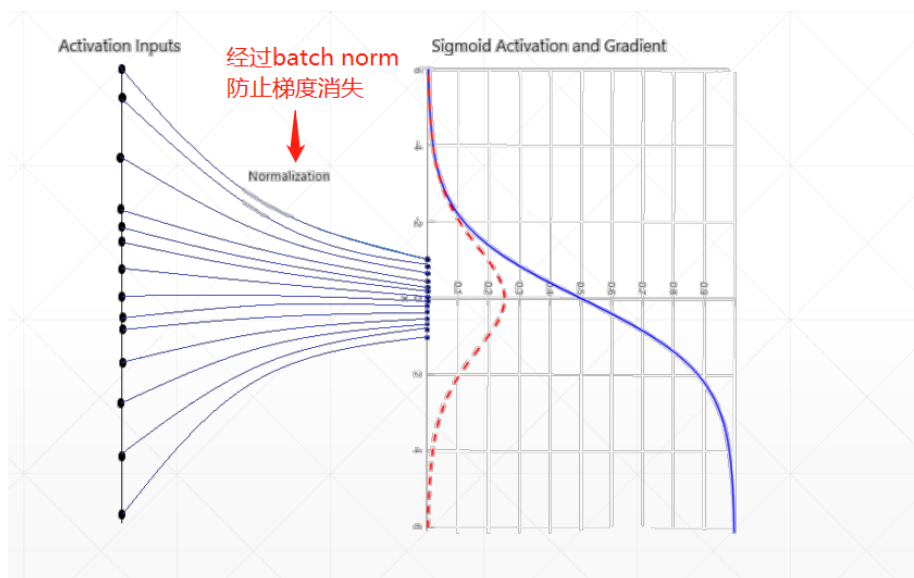


## DenseNet



## Batch Norm


直观的解释:



## Image Normalization

```
def normalize(x, mean, std):
    # x: [b, h, w, c]
    x = x - mean
    x = x / std
    return x
```

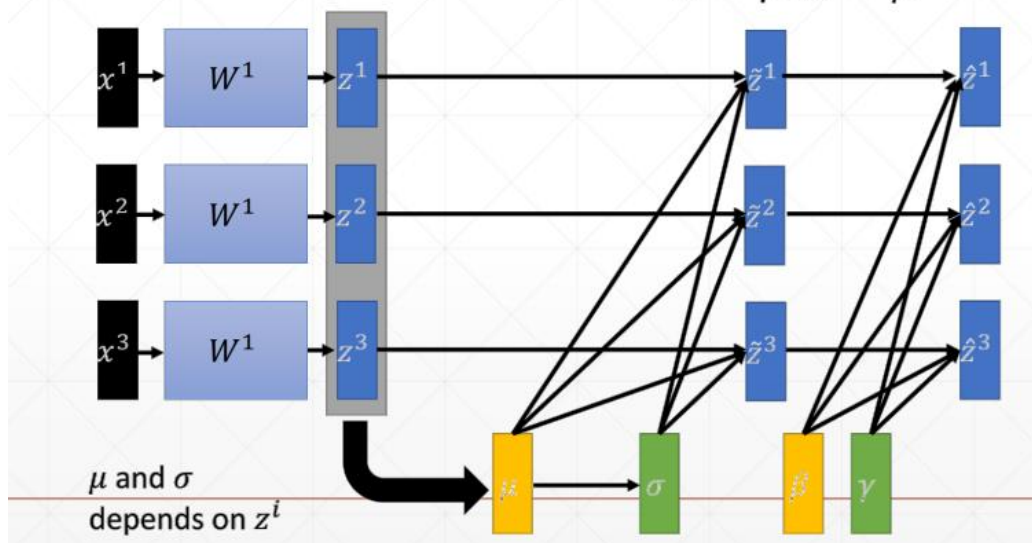
```
(mean=[0.485, 0.456, 0.406],
std=[0.229, 0.224, 0.225])
```



## Batch normalization

$$\tilde{z}^i = \frac{z^i - \mu}{\sigma}$$

gamma 和 beta 通过学习得到  $\hat{z}^i = \gamma \odot \tilde{z}^i + \beta$



## layers.BatchNormization

- `net = layers.BatchNormization()`
  - `axis=-1,`
  - `center=True,`
  - `scale=True`
  - `trainable=True`
- `net(x, training=None)`

```

In [3]: net=layers.BatchNormalization()
In [5]: x=tf.random.normal([2,3])
In [6]: out=net(x)

In [7]: net.trainable_variables
[<tf.Variable 'batch_normalization/gamma:0' shape=(3,) dtype=float32, numpy=array([1., 1., 1.],
dtype=float32)>,
 <tf.Variable 'batch_normalization/beta:0' shape=(3,) dtype=float32, numpy=array([0., 0., 0.],
dtype=float32)>]

In [8]: net.variables
[<tf.Variable 'batch_normalization/gamma:0' shape=(3,) dtype=float32, numpy=array([1., 1., 1.],
dtype=float32)>,
 <tf.Variable 'batch_normalization/beta:0' shape=(3,) dtype=float32, numpy=array([0., 0., 0.],
dtype=float32)>,
 <tf.Variable 'batch_normalization/moving_mean:0' shape=(3,) dtype=float32, numpy=array([0., 0.,
0.], dtype=float32)>,
 <tf.Variable 'batch_normalization/moving_variance:0' shape=(3,) dtype=float32, numpy=array([1.,
1., 1.], dtype=float32)>]

```

## 规范化写法

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_{1\dots m}\}$ ;  
Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

**Algorithm 1:** Batch Normalizing Transform, applied to activation  $x$  over a mini-batch.

```

In [15]: x=tf.random.normal([2,4,4,3],mean=1.,stddev=0.5)
In [16]: net=layers.BatchNormalization(axis=3)
In [17]: out=net(x)

In [18]: net.variables
[<tf.Variable 'batch_normalization_2/gamma:0' shape=(3,) dtype=float32, numpy=array([1., 1.,
1.], dtype=float32)>,
 <tf.Variable 'batch_normalization_2/beta:0' shape=(3,) dtype=float32, numpy=array([0., 0., 0.],
dtype=float32)>,
 <tf.Variable 'batch_normalization_2/moving_mean:0' shape=(3,) dtype=float32, numpy=array([0.,
0., 0.], dtype=float32)>,
 <tf.Variable 'batch_normalization_2/moving_variance:0' shape=(3,) dtype=float32,
numpy=array([1., 1., 1.], dtype=float32)>]

```

```

for i in range(10):
    with tf.GradientTape() as tape:
        out = net(x, training=True)
        loss = tf.reduce_mean(tf.pow(out,2)) - 1
        grads = tape.gradient(loss, net.trainable_variables)
        optimizer.apply_gradients(zip(grads, net.trainable_variables))

backward(10 steps):
[<tf.Variable 'batch_normalization/gamma:0' shape=(3,) dtype=float32, numpy=array([0.93549937,
0.9356556 , 0.9355564 ], dtype=float32)>,
 <tf.Variable 'batch_normalization/beta:0' shape=(3,) dtype=float32, numpy=array([ 1.3411044e-
09,  1.3411045e-08, -4.0978188e-10], dtype=float32)>...]

```

## 反向传播



## 前向传播

```
In [15]: x=tf.random.normal([2,4,4,3],mean=1.,stddev=0.5)
In [16]: net=layers.BatchNormalization(axis=3)
In [19]: out=net(x,training=True)
In [20]: net.variables
[<tf.Variable 'batch_normalization_2/gamma:0' shape=(3,) dtype=float32, numpy=array([1., 1., 1.], dtype=float32)>,
 <tf.Variable 'batch_normalization_2/beta:0' shape=(3,) dtype=float32, numpy=array([0., 0., 0.], dtype=float32)>,
 <tf.Variable 'batch_normalization_2/moving_mean:0' shape=(3,) dtype=float32, numpy=array([0.00992113, 0.00976678, 0.00942467], dtype=float32)>,
 <tf.Variable 'batch_normalization_2/moving_variance:0' shape=(3,) dtype=float32, numpy=array([0.9923687 , 0.9918039 , 0.99186534], dtype=float32)>]

In [21]: for i in range(100):out=net(x,training=True)
In [22]: net.variables
[<tf.Variable 'batch_normalization_2/gamma:0' shape=(3,) dtype=float32, numpy=array([1., 1., 1.], dtype=float32)>,
 <tf.Variable 'batch_normalization_2/beta:0' shape=(3,) dtype=float32, numpy=array([0., 0., 0.], dtype=float32)>,
 <tf.Variable 'batch_normalization_2/moving_mean:0' shape=(3,) dtype=float32, numpy=array([0.63259894, 0.6227575 , 0.60094345], dtype=float32)>,
 <tf.Variable 'batch_normalization_2/moving_variance:0' shape=(3,) dtype=float32, numpy=array([0.5134074, 0.4773918, 0.4813124], dtype=float32)>]
```

## 实战:

```
import os
os.environ['TF_CPP_MIN_LOG_LEVEL'] = '2'
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers, Sequential

class BasicBlock(layers.Layer):

    def __init__(self, filter_num, stride=1):
        super(BasicBlock, self).__init__()

        self.conv1 = layers.Conv2D(filter_num, (3,3), strides=stride, padding='same')
        self.bn1 = layers.BatchNormalization()
        self.relu = layers.Activation('relu')

        self.conv2 = layers.Conv2D(filter_num, (3,3), strides=1, padding='same')
        self.bn2 = layers.BatchNormalization()

        if stride != 1:
            self.downsample = Sequential()
            self.downsample.add(layers.Conv2D(filter_num, (1,1), strides=stride))
        else:
            self.downsample = lambda x:x
```

```

def call(self, inputs, training=None):
    # [b, h, w, c]
    out = self.conv1(inputs)
    out = self.bn1(out)
    out = self.relu(out)

    out = self.conv2(out)
    out = self.bn2(out)

    identity = self.downsample(inputs)

    output = layers.add([out, identity])
    # output = self.relu(output) 等价于下面的
    output = tf.nn.relu(output)

    return output

class ResNet(keras.Model):

    def __init__(self, layer_dims, num_classes=100):
        super(ResNet, self).__init__()

        self.stem = Sequential([layers.Conv2D(64, [3, 3], strides=(1, 1))
                                ,
                                layers.BatchNormalization(),
                                layers.Activation('relu'),
                                layers.MaxPool2D(pool_size=(2, 2), strides=[1, 1], padding='same')])

        self.layer1 = self.build_resblock(64, layer_dims[0])
        self.layer2 = self.build_resblock(128, layer_dims[1], stride=2)
        self.layer3 = self.build_resblock(256, layer_dims[2], stride=2)
        self.layer4 = self.build_resblock(512, layer_dims[3], stride=2)

        self.avgpool = layers.GlobalAveragePooling2D()
        self.fc = layers.Dense(num_classes)

    def call(self, inputs, training=None):
        x = self.stem(inputs)
        x = self.layer1(x)
        x = self.layer2(x)

```

```

        x = self.layer3(x)
        x = self.layer4(x)

        x = self.avgpool(x)
        x = self.fc(x)

        return x

    def build_resblock(self, filter_num, blocks, stride=1):

        res_blocks = Sequential()
        res_blocks.add(BasicBlock(filter_num, stride))

        for _ in range(1, blocks):
            res_blocks.add(BasicBlock(filter_num, stride=1))

        return res_blocks

def resnet18():
    return ResNet([2, 2, 2, 2])

def resnet34():
    return ResNet([3, 4, 6, 3])

```

主代码

```

import os
os.environ['TF_CPP_MIN_LOG_LEVEL'] = '2'
import tensorflow as tf
from tensorflow.keras import layers, optimizers, datasets, Sequential
from resnet import resnet18
tf.random.set_seed(2345)

# 数据预处理与加载数据集
def preprocess(x, y):
    x = tf.cast(x, dtype=tf.float32) / 255.
    y = tf.cast(y, dtype=tf.int32)
    return x, y

(x, y), (x_test, y_test) = datasets.cifar100.load_data()
y = tf.squeeze(y, axis=1)
y_test = tf.squeeze(y_test, axis=1)

train_db = tf.data.Dataset.from_tensor_slices((x, y))
train_db = train_db.shuffle(1000).map(preprocess).batch(128)
test_db = tf.data.Dataset.from_tensor_slices((x, y))
test_db = test_db.map(preprocess).batch(64)

```



```

def main():
    # [b, 32, 32, 3] ==> [b, 1, 1, 512]
    model = resnet18()
    model.build(input_shape=(None, 32, 32, 3))
    optimizer = optimizers.Adam(lr=1e-3)

    for epoch in range(50):
        for step, (x,y) in enumerate(train_db):
            with tf.GradientTape() as tape:
                # [b, 512] ==> [b, 100]
                logits = model(x)
                # 对 y 进行 one_hot 编码
                y_onehot = tf.one_hot(y, depth=100)
                # 计算损失
                loss = tf.losses.categorical_crossentropy(y_onehot, l
ogits, from_logits=True)
                loss = tf.reduce_mean(loss)

                grads = tape.gradient(loss, model.trainable_variables)
                optimizer.apply_gradients(zip(grads, model.trainable_vari
ables))

            if step % 100 == 0:
                print(epoch, step, 'loss', float(loss))

        total_num = 0
        total_correct = 0
        for x, y in test_db:
            logits = model(out)
            prob = tf.nn.softmax(logits, axis=1)
            pred = tf.argmax(prob, axis=1)
            pred = tf.cast(pred, dtype=tf.int32)

            correct = tf.cast(tf.equal(pred, y), dtype=tf.int32)
            correct = tf.reduce_sum(correct)

            total_num += x.shape[0]
            total_correct += int(correct)

        acc = total_correct / total_num
        print(epoch, 'acc:', acc)

main()

```